

Introduction

This notebook contains code for a homework assignment on text classification of news articles into 5 categories. It explores different feature engineering methods, trains neural network models using 5-fold cross validation, and evaluates performance. The best model is then used to predict labels on a test set.

1. Imports and Setup

```
!pip install rake-nltk

Collecting rake-nltk
  Downloading rake_nltk-1.8.6-py3-none-any.whl (9.1 kB)
Requirement already satisfied: nltk<4.0.0,>=3.6.2 in /usr/local/lib/python3.10/dist-packages (from rake-nltk) (3.8.1)
Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from nltk<4.0.0,>=3.6.2->rake-nltk) (8.1.7)
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from nltk<4.0.0,>=3.6.2->rake-nltk) (1.3.2)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.10/dist-packages (from nltk<4.0.0,>=3.6.2->rake-nltk) (2023.12.25)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from nltk<4.0.0,>=3.6.2->rake-nltk) (4.66.2)
Installing collected packages: rake-nltk
Successfully installed rake-nltk-1.8.6

• Import required libraries
• Mount Google Drive
• download nltk packages

from google.colab import drive
drive.mount('/content/drive')
import numpy as np
import random
from tqdm import tqdm
import pandas as pd
import string
import nltk
nltk.download('punkt')
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer

from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.model_selection import kfold

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader

import matplotlib.pyplot as plt
import gensim.downloader as api
import gensim

Mounted at /content/drive
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip
```

Loading training and test data into pandas dataframe

Part 1 Data Loading and Preprocessing

Question 1 Preprocess the raw training data. You can use the code from Homework 1. You are required to construct other features, such n-grams or keyword extractions. (15pt)

- Load train and test CSV files
- Explore data
- Define and apply text preprocessing function

```
train_data_path = '/content/drive/MyDrive/Colab Notebooks/Homework2/24_train_1.csv'
test_data_path = '/content/drive/MyDrive/Colab Notebooks/Homework2/news-test.csv'
train_data = pd.read_csv(train_data_path)
test_data = pd.read_csv(test_data_path)
```

```
print(train_data)
print(train_data.info())
print(train_data['category'])

# Sample data preview
#      ArticleID      Text
# 0      1429  sfa health report over nikolas the scottish...
# 1      1896  parmalat to return to stockmarket parmalat th...
# 2      1859  edu blast arsenal arsenal s brazilian mofiel...
# 3      2178  herman decides to quit david cup the herman ha...
# 4       194  french suitor holds lse meeting european stock...
# ...
# 995     1258  blair damaged by blunkett row a majority of ...
# 996     1639  a november to remember last saturday one new...
# 997       916  highbury tunnel players in clear the football ...
# 998     2217  top stars join us tsunami tv show brad pitt r...
# 999       982  eastwood s baby scoop top oscar clim eastw...
```

```
Category
0      sport
1      business
2      sport
3      sport
4      business
..      ...
995     politics
996     sport
997     sport
998  entertainment
999  entertainment

[1000 rows x 3 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 3 columns):
 #   Column  Non-Null Count  Dtype
---  ---
 0   ArticleID  1000 non-null    int64
 1   Text      1000 non-null    object
 2   Category  1000 non-null    object
dtypes: int64(1), object(2)
memory usage: 23.6+ KB
None
```

Question 1.a: Run Neural Networks with the 2-hidden layers, each has 128 neurons, extracting features by CountVectorizer() as the original features. Use 5-fold cross-validation to evaluate the performance.

This section loads the raw training and test data from CSV files into Pandas dataframes. It then preprocesses the text by:

- Tokenizing
- Lowercasing
- Removing punctuation
- Removing non-alphabetic tokens
- Removing stopwords
- Stemming

The preprocessed text is added as a new column in the dataframes.

It then extracts CountVectorizer features and runs a 2-layer neural network with 5-fold cross validation.

```
# Preprocessing function
def preprocess_text(text):
    # Tokenization
    tokens = nltk.word_tokenize(text)

    # Lowercasing
    tokens = [w.lower() for w in tokens]

    # Removing punctuation
    table = str.maketrans('', '', string.punctuation)
    stripped = [w.translate(table) for w in tokens]

    # Removing non-alphabetic tokens
    words = [word for word in stripped if word.isalpha()]

    # Removing stopwords
    stop_words = set(stopwords.words('english'))
    words = [w for w in words if not w in stop_words]

    # Stemming
    porter = PorterStemmer()
    stemmed = [porter.stem(word) for word in words]

    return ' '.join(stemmed)

# Apply preprocessing to training and test data
train_data['Processed_Text'] = train_data['Text'].apply(preprocess_text)
test_data['Processed_Text'] = test_data['Text'].apply(preprocess_text)
```

Feature Engineering

Question 1.b: Feature exploration. Use other features like TFIDF, or any word embeddings provided by other packages like GloVe with gensim, or BERT. Use 5-fold cross-validation to evaluate the performance of your Neural Network.

- This section generates various features from the preprocessed text:
 - Ngrams (bigrams)
 - TF-IDF vectors
 - GloVe word embeddings (averaging word vectors for each document)
 - Word2Vec embeddings (training a Word2Vec model on the corpus and averaging word vectors for each document)

```
# Create n-grams
def generate_ngrams(text, n_gram=1):
    token = [token for token in text.split(" ") if token != "" if token not in stopwords.words('english')]
    ngrams = zip(*[token[i:] for i in range(n_gram)])
    return " ".join(ngram for ngram in ngrams)

# Generate n-grams for training data
train_data['Processed_Text_ngrams'] = train_data['Processed_Text'].apply(lambda x: generate_ngrams(x, n_gram=2))

# Generate n-grams for test data
test_data['Processed_Text_ngrams'] = test_data['Processed_Text'].apply(lambda x: generate_ngrams(x, n_gram=2))

# Load GloVe embeddings
glove_vectors = api.load("glove-twitter-100")
glove_embeddings = {}

for doc in train_data['Processed_Text_ngrams']:
    doc_embedding = []
    for word in doc:
        if word in glove_vectors:
            doc_embedding.append(glove_vectors[word])
        else:
            doc_embedding = np.mean(doc_embedding, axis=0)
    else:
        doc_embedding = np.zeros(100)
    glove_embeddings.append(doc_embedding)

glove_embeddings = np.array(glove_embeddings)

# Create TF-IDF features
tfidf = TfidfVectorizer(ngram_range=(1,2))
tfidf_train = tfidf.fit_transform(train_data['Processed_Text_ngrams']).apply(lambda x: ' '.join(x))
tfidf_test = tfidf.transform(test_data['Processed_Text_ngrams']).apply(lambda x: ' '.join(x))

# Train Word2Vec model
w2v_model = gensim.models.Word2Vec(train_data['Processed_Text_ngrams'], vector_size=100, window=5, min_count=1, workers=4)
w2v_embeddings = []
for doc in train_data['Processed_Text_ngrams']:
    doc_embedding = []
    for word in doc:
        if word in w2v_model.wv:
            doc_embedding.append(w2v_model.wv[word])
        else:
            doc_embedding = np.mean(doc_embedding, axis=0)
    else:
        doc_embedding = np.zeros(100)
    w2v_embeddings.append(doc_embedding)

w2v_embeddings = np.array(w2v_embeddings)
```

Question 1.c. Describe how you generate features. (5pt)

The features are generated as follows:

- CountVectorizer: Counts the frequency of each word in each document
- N-grams: Generates bigrams from the preprocessed text
- TF-IDF: Computes TF-IDF weights for each word in each document
- GloVe: Averages the pre-trained GloVe embeddings for each word in a document to get a document embedding Word2Vec: Trains a Word2Vec model on the corpus and averages the learned word embeddings for each document

Model Definition

This section defines the architecture of the neural network model (NewsClassifier) using PyTorch. The model has an embedding layer, two hidden layers with ReLU activation, and a softmax output layer. The forward pass of the model is defined in the forward() method.

```
# Define neural network model
class NewsClassifier(nn.Module):
    def __init__(self, embedding_dim, hidden_dim, output_dim):
        super(NewsClassifier, self).__init__()
        self.fc1 = nn.Linear(embedding_dim, hidden_dim)
        self.fc2 = nn.Linear(hidden_dim, hidden_dim)
        self.fc3 = nn.Linear(hidden_dim, output_dim)
        self.relu = nn.ReLU()

    def forward(self, x):
        out = self.fc1(x)
        out = self.relu(out)
        out = self.fc2(out)
        out = self.relu(out)
        out = self.fc3(out)
        return out

# Define training function
def train_model(model, train_loader, criterion, optimizer, device):
    model.train()
    running_loss = 0.0
    running_corrects = 0
    for inputs, labels in train_loader:
        inputs = inputs.to(device)
        labels = labels.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        _, preds = torch.max(outputs, 1)
        running_loss += loss.item() * inputs.size(0)
        running_corrects += torch.sum(preds == labels.data)
    loss.backward()
    optimizer.step()
    epoch_loss = running_loss / len(train_loader.dataset)
    epoch_acc = running_corrects.double() / len(train_loader.dataset)
    print(f"Training Loss: (epoch_loss:4f), Training Accuracy: (epoch_acc:4f)")
    return epoch_loss, epoch_acc

# Define evaluation function
def eval_model(model, test_loader, criterion, device):
    model.eval()
    running_loss = 0.0
    running_corrects = 0
    with torch.no_grad():
        for inputs, labels in test_loader:
            inputs = inputs.to(device)
            labels = labels.to(device)
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            _, preds = torch.max(outputs, 1)
            running_loss += loss.item() * inputs.size(0)
            running_corrects += torch.sum(preds == labels.data)
    epoch_loss = running_loss / len(test_loader.dataset)
    epoch_acc = running_corrects.double() / len(test_loader.dataset)
    print(f"Validation Loss: (epoch_loss:4f), Validation Accuracy: (epoch_acc:4f)")
    return epoch_loss, epoch_acc

# Define dataset class
class NewsDataset(Dataset):
    def __init__(self, features, labels):
        self.features = torch.tensor(features, dtype=torch.float32) # Convert features to float32
        self.labels = torch.tensor(labels, dtype=torch.long) # Labels should be of type long for classification

    def __len__(self):
        return len(self.features)

    def __getitem__(self, idx):
        return self.features[idx], self.labels[idx]

# Define function to run 5-fold cross-validation
def run_cross_validation(model, features, labels, batch_size, num_epochs, device):
    kfold = KFold(n_splits=5, shuffle=True, random_state=42)
    train_accuracies = []
    val_accuracies = []

    for fold, (train_idx, val_idx) in enumerate(kfold.split(features)):
        print(f"Fold {fold+1}")
        train_dataset = NewsDataset(features[train_idx], labels[train_idx])
        val_dataset = NewsDataset(features[val_idx], labels[val_idx])

        train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
        val_loader = DataLoader(val_dataset, batch_size=batch_size)

        model = NewsClassifier(features.shape[1], 128, len(np.unique(labels))).to(device)
        criterion = nn.CrossEntropyLoss()
        optimizer = optim.Adam(model.parameters())

        for epoch in range(num_epochs):
            train_loss, train_acc = train_model(model, train_loader, criterion, optimizer, device)
            val_loss, val_acc = eval_model(model, val_loader, criterion, device)
            print(f"Epoch {epoch+1}/{num_epochs} - Train Loss: (train_loss:4f), Train Acc: (train_acc:4f), Val Loss: (val_loss:4f), Val Acc: (val_acc:4f)")

            train_accuracies.append(train_acc.cpu().numpy())
            val_accuracies.append(val_acc.cpu().numpy())

        return np.mean(train_accuracies), np.std(train_accuracies), np.mean(val_accuracies), np.std(val_accuracies)

# Convert labels to numerical values
label_map = {'sport': 0, 'business': 1, 'politics': 2, 'entertainment': 3, 'tech': 4}
train_labels = train_data['Category'].map(label_map).values
test_labels = test_data['Category'].map(label_map).values
```

Model Training with Cross Validation

- This section defines functions for training (train_model) and evaluating (eval_model) the neural network model.
- It also defines a custom Dataset class (NewsDataset) for loading the features and labels. The run_cross_validation function performs 5-fold cross validation with a given feature set.
- It trains the model on the training set and evaluates on the validation set for each fold.
- It returns the average training and validation accuracies and their standard deviations.
- Cross validation is run with CountVectorizer, TF-IDF, GloVe, and Word2Vec features. The results are printed in a table and visualized in a bar chart.

Double-click (or enter) to edit

3/7


```
# Define the learning rates to explore
learning_rates = [0.0001, 0.0003, 0.001, 0.003, 0.01, 0.03, 0.1]

# Run the cross-validation for each learning rate
lr_results = run_cross_validation_lr(NewClassifier, tfidf_train.toarray(), train_labels, batch_size=32, num_epochs=10, device=torch.device('cuda' if torch.cuda.is_available() else 'cpu'), learning_rates=learning_rates)


# Organize the results in a table
print("Learning Rate\tTrain Accuracy\tTrain Std\tValidation Accuracy\tValidation Std")
for result in lr_results:
    print(f"{result['lr']}\t{result['train_acc_mean']:.3f}\t{result['train_acc_std']:.3f}\t{result['val_acc_mean']:.3f}\t{result['val_acc_std']:.3f}")

Learning Rate: 0.0001
Fold 1
Fold 2
Fold 3
Fold 4
Fold 5
Learning Rate: 0.0003
Fold 1
Fold 2
Fold 3
Fold 4
Fold 5
Learning Rate: 0.001
Fold 1
Fold 2
Fold 3
Fold 4
Fold 5
Learning Rate: 0.003
Fold 1
Fold 2
Fold 3
Fold 4
Fold 5
Learning Rate: 0.01
Fold 1
Fold 2
Fold 3
Fold 4
Fold 5
Learning Rate: 0.03
Fold 1
Fold 2
Fold 3
Fold 4
Fold 5
Learning Rate: 0.1
Fold 1
Fold 2
Fold 3
Fold 4
Fold 5
Learning Rate Train Accuracy Train Std Validation Accuracy Validation Std
0.0001 1.000 0.000 0.952 0.033
0.0003 1.000 0.000 0.976 0.009
0.001 1.000 0.000 0.976 0.009
0.003 1.000 0.000 0.974 0.007
0.01 1.000 0.000 0.958 0.026
0.03 1.000 0.000 0.951 0.014
0.1 0.995 0.003 0.930 0.024

• Visualize accuracy vs learning rate

# Extract learning rates and accuracies for plotting
lrs = [result['lr'] for result in lr_results]
train_accs = [result['train_acc_mean'] for result in lr_results]
val_accs = [result['val_acc_mean'] for result in lr_results]

plt.figure(figsize=(10, 6))
plt.plot(lrs, train_accs, label='Train Accuracy', marker='o')
plt.plot(lrs, val_accs, label='Validation Accuracy', marker='o')
plt.xscale('log')
plt.xlabel('Learning Rate')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy vs Learning Rate')
plt.legend()
plt.show()
```



Optimizer Experiments

Question 2.C: Use 5-fold cross-validation to evaluate the performance w.r.t. optimizers, you could use the feature engineering method that has the best performance from Question 1.

- This section compares the performance of different optimizers (SGD, Adam, RMSprop) using 5-fold cross validation.
- The run_cross_validation_optimizer function trains and evaluates the model with each optimizer.
- The results are printed in a table and visualized in a bar chart of train and validation accuracy for each optimizer.

```
def run_cross_validation_optimizer(model_class, features, labels, batch_size, num_epochs, device, optimizers):
    kfold = KFold(n_splits=5, shuffle=True, random_state=42)
    results = {}

    for opt_name, opt_func in optimizers.items():
        print(f"Optimizer: {opt_name}")
        train_accuracies = []
        val_accuracies = []

        for fold, (train_idx, val_idx) in enumerate(kfold.split(features)):
            print(f"Fold {fold+1}")
            train_dataset = NewDataset(features[train_idx], labels[train_idx])
            val_dataset = NewDataset(features[val_idx], labels[val_idx])

            train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
            val_loader = DataLoader(val_dataset, batch_size=batch_size)

            model = model_class(features.shape[1], 128, len(np.unique(labels))).to(device)
            criterion = nn.CrossEntropyLoss()
            optimizer = opt_func(model.parameters())

            for epoch in range(num_epochs):
                train_loss, train_acc = train_model(model, train_loader, criterion, optimizer, device)
                val_loss, val_acc = eval_model(model, val_loader, criterion, device)

            train_accuracies.append(train_acc.item())
            val_accuracies.append(val_acc.item())

        results.append({
            'optimizer': opt_name,
            'train_acc_mean': np.mean(train_accuracies),
            'train_acc_std': np.std(train_accuracies),
            'val_acc_mean': np.mean(val_accuracies),
            'val_acc_std': np.std(val_accuracies),
        })

    return results

# Define the optimizers to explore
optimizers = {
    'SGD': lambda params: optim.SGD(params, lr=0.001),
    'Adam': lambda params: optim.Adam(params, lr=0.001),
    'RMSprop': lambda params: optim.RMSprop(params, lr=0.001),
}

# Run the cross-validation for each optimizer
opt_results = run_cross_validation_optimizer(NewClassifier, tfidf_train.toarray(), train_labels, batch_size=32, num_epochs=10, device=torch.device('cuda' if torch.cuda.is_available() else 'cpu'), optimizers=optimizers)

# Organize the results in a table
print("Optimizer\tTrain Accuracy\tTrain Std\tValidation Accuracy\tValidation Std")
for result in opt_results:
    print(f"{result['optimizer']}\t{result['train_acc_mean']:.3f}\t{result['train_acc_std']:.3f}\t{result['val_acc_mean']:.3f}\t{result['val_acc_std']:.3f}")
```

```
##### SGD #####
Training Loss: 0.0021, Training Accuracy: 1.0000
Validation Loss: 0.1255, Validation Accuracy: 0.9600
Training Loss: 0.0013, Training Accuracy: 1.0000
Validation Loss: 0.1219, Validation Accuracy: 0.9600
Training Loss: 0.0010, Training Accuracy: 1.0000
Validation Loss: 0.1185, Validation Accuracy: 0.9600
Training Loss: 0.0007, Training Accuracy: 1.0000
Validation Loss: 0.1159, Validation Accuracy: 0.9600
Training Loss: 0.0005, Training Accuracy: 1.0000
Validation Loss: 0.1116, Validation Accuracy: 0.9600
Training Loss: 0.0004, Training Accuracy: 1.0000
Validation Loss: 0.1116, Validation Accuracy: 0.9600
Training Loss: 0.0003, Training Accuracy: 1.0000
Validation Loss: 0.1097, Validation Accuracy: 0.9600
Fold 3
Training Loss: 0.7928, Training Accuracy: 0.7838
Validation Loss: 0.1516, Validation Accuracy: 0.9900
Training Loss: 0.0212, Training Accuracy: 0.9900
Validation Loss: 0.1063, Validation Accuracy: 0.9900
Training Loss: 0.0045, Training Accuracy: 1.0000
Validation Loss: 0.0934, Validation Accuracy: 0.9900
Training Loss: 0.0025, Training Accuracy: 1.0000
Validation Loss: 0.0862, Validation Accuracy: 0.9950
Training Loss: 0.0015, Training Accuracy: 1.0000
Validation Loss: 0.0811, Validation Accuracy: 0.9850
Training Loss: 0.0010, Training Accuracy: 1.0000
Validation Loss: 0.0773, Validation Accuracy: 0.9850
Training Loss: 0.0007, Training Accuracy: 1.0000


# Extract optimizer names and accuracies for plotting
opt_names = [result['optimizer'] for result in opt_results]
train_accs = [result['train_acc_mean'] for result in opt_results]
val_accs = [result['val_acc_mean'] for result in opt_results]

x = np.arange(len(opt_names))
width = 0.35

fig, ax = plt.subplots(figsize=(8, 6))
ax.bar(x = width/2, train_accs, width, label='Train Accuracy')
ax.bar(x = width/2, val_accs, width, label='Validation Accuracy')

ax.set_ylabel('Accuracy')
ax.set_title('Training and Validation Accuracy for Different Optimizers')
ax.set_xticks(x)
ax.set_xticklabels(opt_names)
ax.legend()

plt.tight_layout()
plt.show()
```



PART 3 : Predict the labels for the testing data (using raw training data and raw testing data). (60pt)

Final Model Training and Test Set Prediction

This section trains the final model on the full training set using the best performing feature set (CountVectorizer) and hyperparameters. It then generates predictions on the test set by:

- Preprocessing the test text data
- Extracting CountVectorizer features
- Loading the features into a test Dataset and DataLoader
- Making predictions with the trained model
- Converting the predicted class indices back to labels
- Printing the article ID and predicted label for each test article

Question 3.a: Describe how you pre-process the data to generate features. (5pt)

The test data is preprocessed using the same steps as the training data: Tokenizing, lowercasing, removing punctuation, non-alphabetic tokens and stopwords, stemming Extracting CountVectorizer features (best performing from Q1)

Question 3.b: Describe how you choose the model and parameters. (5pt)

The final model is a 2-layer neural network with 128 hidden units each, using CountVectorizer features. The key hyperparameters (learning rate, optimizer) are chosen based on the cross validation results from Q2.

Question 3.c: Describe the performance of your chosen model and parameter on the training data. (5pt)

The performance of the final model on the full training set is reported, including the training loss and accuracy for each epoch.

Question 3.d: The final classification models to be used in this question are limited to random forest, neural networks, and ensemble methods. It is OK to use other models to do feature engineering. (45pt)

The trained neural network model is used to generate predictions on the test set. The test data is preprocessed, features are extracted, and the trained model is applied to get predicted class probabilities. The class with the highest probability is taken as the predicted label for each test document. The article ID and predicted label are printed out for

```
# Convert labels to numerical values
label_map = {'sport': 0, 'business': 1, 'politics': 2, 'entertainment': 3, 'tech': 4}
train_labels = train_data['Category'].map(label_map).values

# Run cross-validation with CountVectorizer features
vectorizer = CountVectorizer()
count_train = vectorizer.fit_transform(train_data['Processed_Text']).toarray()
count_test = vectorizer.transform(test_data['Processed_Text']).toarray()

# Initialize the neural network model
model = nn.LstmClassifier(count_train.shape[1], 128, len(np.unique(train_labels)))
device = torch.device('cpu' if torch.cuda.is_available() else 'cpu')
model.to(device)

# Define the loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters())

# Create a dataset and loader for the training data
train_dataset = nn.DataLoader(count_train, train_labels)
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)

# Train the model
num_epochs = 10
for epoch in range(num_epochs):
    train_loss, train_acc = train_model(model, train_loader, criterion, optimizer, device)
    print(f"Epoch {epoch+1}/{num_epochs} - Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f}")

# Predict the labels for the testing data
test_dataset = nn.DataLoader(count_test, np.zeros(len(count_test))) # Dummy labels for the test set
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)
predictions = []
model.eval()
with torch.no_grad():
    for inputs, _ in test_loader:
        inputs = inputs.to(device)
        outputs = model(inputs)
        _, preds = torch.max(outputs, 1)
        predictions.extend(preds.cpu().numpy())

# Convert numerical predictions back to category names
inverse_label_map = {v: k for k, v in label_map.items()}
predicted_categories = [inverse_label_map[pred] for pred in predictions]

# Print the article ID and the predicted categories
for article_id, category in zip(test_data['ArticleId'], predicted_categories):
    print(f"Article ID: {article_id}, Predicted Category: {category}")

Training Loss: 0.7579, Training Accuracy: 0.8520
Epoch 1/10 - Train Loss: 0.7579, Train Acc: 0.8520
Training Loss: 0.0209, Training Accuracy: 0.9900
Epoch 2/10 - Train Loss: 0.0209, Train Acc: 0.9900
Training Loss: 0.0026, Training Accuracy: 1.0000
Epoch 3/10 - Train Loss: 0.0026, Train Acc: 1.0000
Training Loss: 0.0014, Training Accuracy: 1.0000
Epoch 4/10 - Train Loss: 0.0014, Train Acc: 1.0000
Training Loss: 0.0007, Training Accuracy: 1.0000
Epoch 5/10 - Train Loss: 0.0007, Train Acc: 1.0000
Training Loss: 0.0005, Training Accuracy: 1.0000
Epoch 6/10 - Train Loss: 0.0005, Train Acc: 1.0000
Training Loss: 0.0004, Training Accuracy: 1.0000
Epoch 7/10 - Train Loss: 0.0004, Train Acc: 1.0000
Training Loss: 0.0003, Training Accuracy: 1.0000
Epoch 8/10 - Train Loss: 0.0003, Train Acc: 1.0000
Training Loss: 0.0003, Training Accuracy: 1.0000
Epoch 9/10 - Train Loss: 0.0003, Train Acc: 1.0000
Training Loss: 0.0002, Training Accuracy: 1.0000
Epoch 10/10 - Train Loss: 0.0002, Train Acc: 1.0000
Article ID: 1018, Predicted Category: sport
Article ID: 1319, Predicted Category: tech
Article ID: 1138, Predicted Category: sport
Article ID: 459, Predicted Category: business
Article ID: 1020, Predicted Category: sport
Article ID: 51, Predicted Category: sport
```

```
Article ID: 2825, Predicted Category: politics
Article ID: 1479, Predicted Category: politics
Article ID: 27, Predicted Category: entertainment
Article ID: 897, Predicted Category: business
Article ID: 1644, Predicted Category: business
Article ID: 563, Predicted Category: tech
Article ID: 765, Predicted Category: politics
Article ID: 2114, Predicted Category: tech
Article ID: 297, Predicted Category: entertainment
Article ID: 1712, Predicted Category: sport
Article ID: 1631, Predicted Category: politics
Article ID: 942, Predicted Category: tech
Article ID: 1549, Predicted Category: entertainment
Article ID: 516, Predicted Category: entertainment
Article ID: 2215, Predicted Category: business
Article ID: 531, Predicted Category: politics
Article ID: 1541, Predicted Category: sport
Article ID: 1348, Predicted Category: business
Article ID: 56, Predicted Category: politics
Article ID: 3118, Predicted Category: sport
Article ID: 338, Predicted Category: business
Article ID: 113, Predicted Category: sport
Article ID: 111, Predicted Category: sport
Article ID: 521, Predicted Category: business
Article ID: 1777, Predicted Category: politics
Article ID: 185, Predicted Category: tech
Article ID: 2185, Predicted Category: business
Article ID: 1550, Predicted Category: business
Article ID: 659, Predicted Category: sport
Article ID: 1837, Predicted Category: sport
Article ID: 59, Predicted Category: sport
Article ID: 383, Predicted Category: business
```